

Interplanetary Consensus (IPC)

Consensus Lab

Abstract

1 Introduction

A blockchain system is a platform for hosting replicated applications (a.k.a. smart contracts in Ethereum [??] or actors in Filecoin [??]). A single system can, at the same time, host many such applications, each of which containing logic for processing inputs (a.k.a. transactions, requests, or messages) and updating its internal state accordingly. The blockchain system stores multiple copies of those applications' state and executes the associated logic. In practice, applications are largely (or even completely) independent. This means that the execution of one application's transactions rarely (or even never) requires accessing the state of another application.

Nevertheless, most of today's blockchain systems process all transactions for all hosted applications (at least logically) sequentially. The whole system maintains a single totally ordered transaction log containing an interleaving of the transactions associated with all hosted applications. The total transaction throughput the blockchain system can handle thus must be shared by all applications, even completely independent ones. This may greatly impair the performance of such a system at scale (in terms of the number of applications). Moreover, if processing a transaction incurs a cost (transaction fee) for the user submitting it, using the system tends to become more expensive when the system is saturated.

The typical application hosted by blockchain systems is asset transfer between users (wallets). It is true that many other applications are often involved in transferring assets and asset transfer may create system-wide dependencies between different parts of the system state. In general, if users interacted in an arbitrary manner (or even uniformly at random), this would indeed be the case. However, in practical systems, users typically tend to cluster in a way that users inside a cluster interact more frequently than users from different clusters. While this "locality" makes it unnecessary to totally order transactions confined to different clusters (in practice, the vast majority of them), many current blockchain systems spend valuable resources on doing so anyway.

An additional issue of such systems is the lack of flexibility in catering for the different hosted applications. Different applications may prefer vastly different trade-offs (in terms of latency, throughput, security, durability, etc...). For example, a high-level money settlement application may require the highest levels of security and durability, but may more easily compromise on performance in terms of transaction latency and throughput. On the other hand, one can imagine a distributed online chess platform (especially one supporting fast chess variants), where most of its state is ephemeral (until the end of the game), but requires high throughput (for many concurrent games) and low latency (few people like waiting 10 minutes for the opponent's move). While the former is an ideal use case for the Bitcoin network, the latter would probably benefit more from being deployed in a single data center.

In the above example, one can also easily imagine those two applications being mostly, but not completely independent. E.g., a chess player may be able to win some money in a chess tournament and later use it to buy some goods outside of the scope of the chess platform. In such a case, few transactions involve both applications (e.g., paying the tournament registration fee and withdrawing the prize money). The rest (e.g., the individual chess moves) are confined to the chess application and can thus be performed much faster and much cheaper (imagine playing chess by posting each move on Bitcoin for comparison).

Interplanetary Consensus (IPC) is a system that enables the deployment of heterogeneous applications on heterogeneous underlying SMR/blockchain platforms, while still allowing them to interact in a secure way. The basic idea behind IPC is dynamically deploying separate, loosely coupled SMR/blockchain systems (that we call *subnets*) to host different (sets of) applications. Each subnet runs its own consensus protocol and maintains its own ordered transaction log.

IPC is organized in a hierarchical fashion, where each subnet, except for one that we call the *rootnet*, is associated with exactly one other subnet (called its *parent*). Conversely, one parent can have arbitrarily many subnets (called *children*) associated with it.

This tree of subnets expresses a hierarchy of trust. All components of a subnet and all users using it are assumed to fully trust their parent and regard it as the ultimate source of truth. Note that, in general, trust in all components of the parent subnet is not required, but the parent system as a whole is always assumed to be correct (for some subnet-specific definition of correctness) by its child.

To facilitate the interaction between different subnets, IPC provides mechanisms for communication between them. In particular:

1. Assuming a common notion of money (assets / tokens / ...) between a parent and a child and accounts that can hold them, IPC also defines

how money is moved between accounts in different subnets.

2. We define the notion of a *checkpoint* as a snapshot of the state of a subnet after having processed a certain sequence of transactions. IPC enables child subnets to place references to their checkpoints inside the state of their parents.

The operating model described above is simple but powerful. In particular, it enables

- Scaling, by using multiple blockchain/SMR platforms to host a large number of applications.
- Optimization of blockchain platforms for applications running on top of them.
- Governance of a child subnet by its parent, by way of the parent serving as the source of truth for the child (and, for example, maintaining the child’s configuration, validator set, and other subnet-specific data).
- “Inheriting” by the subnet of some of its parent’s security and trustworthiness, by periodically anchoring its state (through checkpoints) in the state of the parent.

In the rest of this document, we describe IPC in detail. In section 2 we define the system model and introduce the necessary terminology. Section 3 describes the main components of IPC and their interfaces. [**TODO:** Finish this when sections are written.]

2 Model

The vocabulary used throughout this document is summarized in Appendix A. The reader is encouraged to read Appendix A before continuing. [**TODO:** Go through and address remaining Marko’s comments of this section]

2.1 Computation and failure model

We model IPC as a distributed (“message-passing”) system consisting of *processes* that communicate by exchanging *messages*¹ over a network. In practice, a process is a program running on a computer, having some state, and reacting to external events and messages received over a communication network. We describe processes as exemplified in Algorithm 1.

¹Network messages are not to be confused with Filecoin actor messages, that this document refers to as transactions.

Algorithm 1: Process definition.

```
1 variable = initial value
2 variable = initial value
3 ...
4 ► process:
5   upon event(params...) do
6     | // Logic to execute atomically
7   upon event(params...) do
8     | // Logic to execute atomically
9   ...
```

A process that performs all the steps exactly as prescribed by the protocols it is participating in is *correct*. A process that stops performing any steps (i.e., *crashes*) or that deviates from the prescribed protocols in any way is *faulty*. If a process is correct or may only fail by crashing, it is *benign*. A non-benign process is *malicious*. [mp: We can remove terms we end up not using...]

In general, faulty processes can be malicious (Byzantine), i.e., we do not put any restrictions on their behavior, except being computationally bounded and thus not being able to subvert standard cryptographic primitives, such as forging signatures or inverting secure hash functions. If the implementation of some component in our design requires additional assumptions on the behavior of faulty processes, they will be stated explicitly.

We use the term *participant* to describe an entity participating in the system that controls one or more processes. All processes controlled by one participant are assumed to be in the same trust domain – they trust each other, i.e., assume each other’s correctness. For example, a participant in the child subnet will probably run multiple processes: one for participating in the child subnet’s protocol (child replica), one for participating in the parent subnet (parent replica), and one process that processes the information from the above two and submits transactions accordingly (IPC agent). We precisely define the replicas and the IPC agent (all of them being processes) in Sections 2.2 and 3. The IPC agent of a participant always assumes that the information it receives from “its own” child replica is correct. However, messages received from another participant’s replica or IPC agent are seen as potentially malicious.

The synchrony assumptions may vary between different components of IPC. We thus state those assumptions whenever necessary, when describing concrete implementations of IPC components.

2.2 State machine replication (SMR) and smart contracts

SMR and replicated state. A *state machine replication (SMR) system*² is a system consisting of processes called *replicas*, each of which locally stores

²In this document, we use the terms “SMR system” and “blockchain” interchangeably.

a copy of (or at least has access to) *replicated state* that it updates over time by applying a sequence of *transactions* to it. Without specifying the details of it, we assume that any process can *submit* a transaction to an SMR system (we call such a process an *SMR client*) and that this transaction will eventually be ordered and applied to the replicated state. We call an SMR system that is part of IPC a *subnet*.

An SMR system guarantees to each correct replica that, after applying n transactions to its local copy of the replicated state, the latter will be identical to any other correct replica's copy of the replicated state after applying n transactions. The replicas achieve this by executing an *ordering protocol* to agree on a common sequence of transactions to apply to the replicated state.

Note that replicas do not necessarily all hold the same replicated state at any instant of real time, since each replica might be processing transactions at a different time. In this context, there is no such thing as “the current replicated state of the SMR system”. There is only the current replicated state of a single replica. The replicated state of the system is only an abstract, logical construct useful for reasoning about transitions from one replicated state to another, happening at individual replicas by applying transactions (at different real times). When referring to a “current” replicated state, we mean the state resulting from the application of a certain number of transactions to the initial state.

Smart contracts. The replicated state of an SMR system can be logically subdivided into multiple *smart contracts* (a.k.a. actors in Filecoin). A smart contract is a portion of the replicated state with well-defined semantics. It defines the logic (e.g., expressed in a programming language, like Solidity in Ethereum) that a replica needs to execute when applying transactions and the new state that results from it.

We model a smart contract as a logical object in the replicated state that contains arbitrary variables representing its state. Its associated logic reacts to *events* triggered by (1) the application of transactions or (2) execution of other (or even own) smart contract logic. We describe smart contracts as exemplified in Algorithm 2.

Algorithm 2: smart contract definition

```

1 variable = initial value
2 variable = initial value
3 ...
4 ► smart contract name:
5   | Function(params...)
   | | // Logic to execute
6   | Function(params...)
   | | // Logic to execute

```

Note that a process usually represent OS-level processes running on some physical machine executing a program, smart contracts are an abstraction over the replicated state of an SMR system and their logic is being executed by all its replicas. While a process can submit a transaction to an SMR system, a smart contract cannot.

Naming. We assign each subnet a name that is unique among all the children of the same parent. Similarly to the notation used for absolute paths in a file system, the name of a child subnet is always prefixed by the name of its parent. For example, subnets P/C and P/D would both be children of subnet P .

Interaction between subnets. In IPC, whole subnets need to interact, i.e., the replicated state of one subnet must react to (changes in) the replicated state of another subnet. As the replicated state of every subnet is distributed among its replicas and evolves independently of other subnets, we must establish a mechanism for interactions between the states of subnets. In particular, we must explicitly link the two replicated states of two subnets. More precisely, for any interaction between two subnets (A and B), define block heights h_A and h_B , such that A 's replicated state at height h_A considers B 's replicated state to have evolved exactly until h_B .

Proofs of Finality. To enable interaction between subnets, we define a *Proof of Finality (PoF)* to be data that proves that an SMR system definitively reached a certain replicated state. Regardless of the SMR system's ordering protocol's approach to finality (e.g., immediate finality for classic BFT protocols, or probabilistic finality in PoW-based systems), a PoF convinces the the proof's verifier that the replicated state the PoF refers to will not be rolled back. For example, for a BFT-based SMR system, a quorum of signatures produced by its replicas can constitute a PoF. We denote by *PoF* (tx) the proof that an SMR system reached a state in which transaction tx already has been applied.

2.3 Representing value

For each pair of subnets in a parent-child relationship, we assume that there exists a notion of *value* (measured in *coins*) common to both subnets.³ Each end user of the SMR system is assumed to have a personal wallet and a corresponding account in some subnet.

We also assume that the submission, ordering, and applications of transactions is associated with a variable cost. Each SMR client submitting a

³One can easily generalize the design to decouple the use of value between a parent and its child, but we stick with using the same kind of value in both subnets for simplicity.

transaction to a subnet is assumed to have an account in that subnet, from which this cost is deducted. If the funds are insufficient, the SMR system ignores the transaction.

3 Components and their Interfaces

We now focus on the interaction between two subnets in a parent-child relation. This interaction comprises running the subnets, observing each other's replicated state, constructing Proofs of Finality, submitting the necessary transactions, and modifying the replicated state accordingly. To enable this interaction, IPC consists of several components and interfaces between them, which we illustrate in Figure 1.

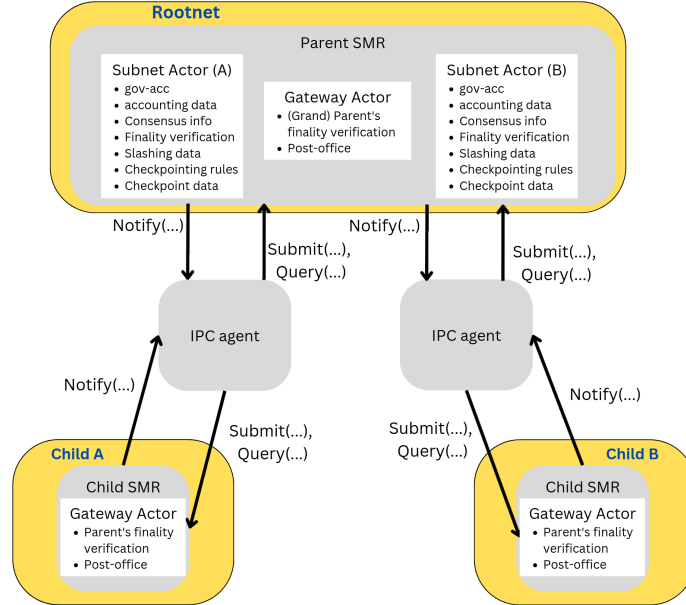


Figure 1: The basic IPC components and their interfaces in an example with one parent and 2 child subnets (A and B).

3.1 Components

IPC components consist of three types of *processes* and two types of *actors* (smart contracts).

3.1.1 Processes

1. **Parent replica:** The process that executes the SMR protocol of the parent subnet. It has a copy of the parent's replicated state, partici-

pates in receiving and ordering transactions and updates the replicated state (including the *SA* smart contract) accordingly.

2. **Child replica:** The process that executes the SMR protocol of the child subnet. It has a copy of the child’s replicated state, participates in receiving and ordering transactions and updates the replicated state (including the *GWA* smart contract) accordingly.
3. **IPC agent:** The process that mediates the interactions between the two subnets. It has access to the replicated states of both the parent and the child (e.g., by sharing a trust domain with a child and a parent replica, or by downloading the replicated state from other replicas) and acts as an SMR client (i.e., submits transactions) to both subnets. It is also responsible for constructing Proofs of Finality (which might involve communication with other processes).

3.1.2 Actors (smart contracts)

1. **Subnet actor (*SA*):** The smart contract in the parent subnet’s replicated state that stores all information about the child subnet that the parent subnet needs. The *SA* is created by the *GWA* (see below) and its functions are invoked by transactions the IPC agent. The state of the *SA* includes:
 - *Accounting data.* This data describes the value that has been deposited to the child. It is considered locked inside the *SA* until it is withdrawn from the child. This data might consist of just a single value representing the sum of all such coins (“aggregated accounts” approach), but might also contain finer-grained information about balances for each account in the child subnet (“segregated accounts” approach).
 - *Ordering protocol (subnet consensus) data.* This is the data (or a pointer to it) that is needed to run the ordering protocol of the child subnet. It is ordering protocol-specific, but is generally expected to contain information such as
 - The ordering protocol used by the subnet.
 - Subnet configuration such as the validator set, voting rights, and collateral deposits.
 - Subnet governance mechanisms, e.g., transaction fees, block rewards, conditions for participation.

[**TODO:** Mention our reference implementation as a concrete example here, saying what information it stores.]

- *Child state finality verification.* Logic to verify that a given child’s replicated state/*tx* is final. or that a particular *tx* has been definitively included in the child’s state. We expect that this logic will

verify a PoF submitted (through transactions) by one or more IPC agents to the *SA*.

- *Slashing*.
 - List of slashable misbehaviors and a proving methodology. That is, for each slashable misbehavior there is a definition of what constitutes a valid proof of misbehavior (*PoM*).
 - Penalties for misbehavior and rewards for reporting *PoM*, as well as the logic performing the actions necessary for slashing in the parent subnet.
- *Checkpointing*.
 - Child subnet’s checkpoint data or a pointer to it.
 - Checkpoint validity rules (and logic enforcing them). E.g., “Checkpoints must be at least every Δ subnet blocks apart”, or “A checkpoint must contain a hash of the previous checkpoint”.

2. **IPC coordinator/gateway actor (*GWA*)**: a smart contract that exists in every subnet in the IPC hierarchy and contains all information and logic the subnet itself needs to hold in order to be part of IPC. The state of the *GWA* includes:

- *Membership data* defining the set of replicas running the *GWA*’s subnet. This may include the identities of the replicas, their network addresses, weights of their votes in the executed consensus protocol (e.g., storage power table, or stake power table), and other subnet-specific membership information.
- *Parent state finality verification*. Analogously to the *SA*’s child state finality verification logic, this is the logic to verify that a state/*tx* is final in the parent subnet, using a PoF submitted as transaction(s) to the child subnet by the IPC agent(s).
- *Inter-subnet transactions service* (denoted postBox). The *GWA* contains a registry of subnets and a functionality that can be used to transfer data from one subnet to another. The postBox specifies the methods and the state locations that are used for these services. This functionality is required for the communication of two smart contracts across subnets.⁴

⁴When inter-subnet data transfer happens between users, they can actively participate in the propagation by submitting transactions to the parent and child subnets. Smart contracts, on the other hand, do not have that power and, therefore, cannot communicate inter-subnets as efficiently as users (*EOA*).

3.2 Interfaces

We now describe the interfaces of the Gateway Actor (*GWA*) and the Subnet Actor (*SA*), by listing the methods that can be invoked through transactions submitted to their respective subnets, as well as that of the IPC Agent process, by listing the events it reacts to. The other two processes, the parent and child replica, interact with the IPC Agent by having the IPC Agent observe the relevant parts of parent's and child's replicated state. The parent and child replicas do not perform any IPC-specific tasks themselves, except for providing the data necessary for IPC Agents to construct Proofs of Finality.

Notation. We refer to an account a in the replicated state of subnet S as $S.a$. To denote a Function a of Smart Contract in the replicated state of a Subnet, we write $Subnet.SmartContract.Function$. E.g., the *GWA*'s function *CreateChild* in subnet P is denoted $P.GWA.CreateChild$. We also use this notation for a transaction tx submitted to subnet P that invokes the function, e.g., $tx = P.GWA.CreateChild(P/C, params)$.

3.2.1 Gateway Actor (*GWA*)

Algorithm 3: *GWA* interface

```

1  ► GWA:
2  |   CreateChild(name, params)
3  |   |   Creates a new SA with the given name and subnet-specific parameters
3  |   |   (such as initial membership, etc.). The subnet governed by the created
3  |   |   SA will be considered the child of the subnet of this GWA.
4  |   RemoveChild(name)
5  |   |   [mp: Is it meaningful to have this functionality at all?]
6  |   Joined(identity, metadata, PoF)
7  |   |   For subnets with explicit membership defined by the SA in their parent,
7  |   |   if PoF is valid, updates the membership data of the GWA to include
7  |   |   the replica with identity and the associated metadata. A valid PoF
7  |   |   means that SA.Join(identity, ..., ...) has been successfully invoked in
7  |   |   the parent's replicated state.
8  |   Leave(identity)
9  |   |   For subnets with explicit membership defined by the SA in their parent,
9  |   |   removes the replica with identity from this subnet's membership.
10 |   Deposited(amt, dest, PoF)
11 |   |   If PoF is valid, adds amt newly minted coins to account dest. A valid
11 |   |   PoF means that a corresponding SA.Deposit(..., amt, dest) has been
11 |   |   successfully invoked in the parent's replicated state.
12 |   Withdraw(src, amt, dest)
13 |   |   Burns amt coins from account src to be returned to the dest account in
13 |   |   the parent subnet.
14 |   Propagate(tx)
15 |   |   Adds the cross-net transaction tx to the list of transactions to be
15 |   |   submitted to the another subnet. (The IPC agents observing the state
15 |   |   of this subnet will pick it up from here and do the actual submission.)

```

3.2.2 Subnet Actor (SA)

Algorithm 4: *SA* interface (governing subnet *C*)

```

1  ► SA:
2  | Join(identity, src, collateral)
3  |   For subnets with explicit membership defined by the SA, adds the
   |   replica with the given identity to the replica set of the subnet governed
   |   by this SA. Join also locks collateral coins in the src account the will
   |   only be released when the replica leaves.
4  | Left(identity, PoF)
5  |   For subnets with explicit membership defined by the SA, if PoF is
   |   valid, adds the replica with the given identity to the replica set of the
   |   subnet governed by this SA. The PoF proves that the corresponding
   |   C.GWA.Leave(identity) has been successfully invoked and the replica
   |   with the given identity is not running subnet C any more.
6  | Deposit(src, amt, dest)
7  |   Locks amt coins on account src to be deposited to the dest account on
   |   the child subnet.
8  | Withdrawn(amt, dest, PoF)
9  |   If PoF is valid (meaning that a corresponding C.GWA.Withdraw(...,
   |   amt, dest) has been successfully invoked), unlocks amt coins on the
   |   dest account.
10 | Checkpoint(chkp, PoF)
11 |   If PoF is valid (meaning that a corresponding child subnet indeed
   |   reached finality on the state represented by chkp), saves chkp in the
   |   replicated state as part of the SA. This will be the most recent state
   |   that the child subnet is considered to have reached.

```

3.2.3 IPC Agent

We assume that an IPC Agent is only responsible for a single pair of parent-child subnets, the state of which it has access to. It reacts to changes in those states triggered by the corresponding invocations of smart contracts, as listed below.

Algorithm 5: IPC Agent interface

```
1 ▶ IPC Agent:
2   upon parent.SA.Join(identity, src, collateral) do
3     Constructs a PoF proving that the invocation of parent.SA.join(identity,
      src, collateral) has been finalized in the parent's replicated state, as
      well as subnet-specific metadata based on the identity of the joining
      replica and the associated collateral and notifies the child subnet by
      submitting a transaction that invokes child.GW.Joied(identity,
      metadata, PoF).
4   upon child.GW.Leave(identity) do
5     Constructs a PoF proving that the invocation of
      child.GW.Leave(identity) has been finalized in the child's replicated
      state and notifies the parent subnet by submitting a transaction that
      invokes parent.SA.Left(identity, PoF)
6   upon parent.SA.Deposit(..., amt, dest) do
7     Constructs a PoF proving that the invocation of parent.SA.Deposit(src,
      amt, dest) has been finalized in the parent's replicated state and
      notifies the child subnet by submitting a transaction that invokes
      child.GW.Deposited(amt, dest, PoF)
8   upon child.GW.Withdraw(..., amt, dest) do
9     Constructs a PoF proving that the invocation of child.GW.Withdraw(...,
      amt, dest) has been finalized in the child's replicated state and notifies
      the parent subnet by submitting a transaction that invokes
      parent.SA.Withdrawn(amt, dest, PoF)
10  upon child.GW.CrossNetTX(tx, destName) do
11    If destName points up the hierarchy, submits tx to the parent subnet.
12  upon parent.GW.CrossNetTX(tx, destName) do
13    If destName points down the hierarchy, submits tx to the child subnet.
```

Note the function pairs *Joined/Leave* of the *GWA* actor (Algorithm 3) and *Join/Left* of the *SA* (Algorithm 4). This is because the intended invocation pattern for several functionalities is as follows (to be detailed in Section 4).

1. The initiating subnet invokes a function on a smart contract (e.g., *SA.Join*).
2. IPC agent notices the invocation, constructs the required PoF, and submits a transaction to the other subnet (e.g., *GWA.Joined*)

4 IPC functionality

IPC exposes the following functionalities:

- Creating child subnets.
- Removing child subnets.
- Depositing coins from an account in a subnet to an account in its child.

- Withdrawing coins from an account in a subnet to an account in its parent.
- Checkpointing - including a checkpoint of a subnet's replicated state in the replicated state of its parent.
- Propagating cross-net-transactions - invoking smart contracts in a subnet through changes in the replicated state of another subnet.

In the following, we describe each functionality in detail.

4.1 Creating a child subnet

Any user of a subnet P can create a new subnet P/C by submitting a transaction $P.GWA.CreateChild(P/C, params)$. This results in the creation of a new subnet actor SA_C in P governing the subnet P/C . The $params$ value describes all the subnet-specific parameters required to initialize the state of $P.SA_C$, such as the initial membership data, the consensus protocol to use, etc.

4.2 Deposits

A deposit is a transfer of funds (of some amount amt) from an account src in the parent subnet P to an account $dest$ in the child subnet P/C . We assume that the owner of src is either running their own IPC Agent to perform the necessary operations described below, or uses another trusted IPC agent to act on their behalf. The deposit is performed as follows:

1. The owner of src submits a transaction $tx = P.SA.Deposit(src, amt, dest)$.
2. The parent subnet orders and executes the *Deposit* transaction (provided src has enough funds) by transferring amt from src to the SA (concretely, to $dest$ account representation within the SA). This effectively locks the funds within the SA smart contract, until the SA smart contract transfers it back to src during a withdrawal (see Section 4.3).
3. When the parent's replicated state that includes tx becomes final (for some SMR-system-specific definition of finality), The IPC agent constructs a $PoF(tx)$ ⁵
4. xThe IPC Agent submits a transaction $tx' = P/C.Deposited(amt, dest, PoF)$ to the child SMR system.
5. Upon ordering tx' , the replicated logic of the child SMR system mints amt new coins and adds them to $dest$.

⁵The exact content of PoF for the transaction tx depends on the implementations of the SMR systems. It might contain, for example, a quorum of replica signatures, a Merkle proof of inclusion, or even be empty.

The events being produced and consumed by the deposit functionality and in Algorithm 6 the pseudocode per component to implement the functionality.

Algorithm 6: Deposit operation

```

1 ► Owner of src:
2   └ submit  $tx = P.SA.Deposit(src, amt, dest)$  to parent subnet
3 ►  $P.SA.Deposit(src, amt, dest)$ :
4   └ move  $amt$  from  $src$  to  $P.SA.accounts.dest$  // "lock" at parent
5 ► IPC agent:
6   └ upon  $tx = P.SA.Deposit$  final at parent do
7     └ create  $PoF$  that  $tx$  is final at parent subnet // see section 7
8     └ submit  $P/C.GWA.Deposited(amt, dest, PoF)$ 
9 ►  $P/C.GWA.Deposited(amt, dest, PoF)$ :
10  └ verify( $PoF$ )
11  └ increase  $dest$  account by  $amt$ 

```

4.3 Withdrawals

A withdrawal is a transfer of funds from an account src in the child subnet P/C to an account $dest$ in the parent subnet P . The *Withdraw* is performed analogously to the *Deposit*, but starting at the child subnet P/C :

1. The owner of src submits a transaction $tx = P/C.GWA.Withdraw(src, amt, dest)$.
2. The child subnet orders and executes the *Withdraw* transaction, burning amt funds in src (provided src has enough funds).
3. When the child's replicated state that includes the transaction becomes final (for some SMR-system-specific definition of finality that has been defined in the SA), the IPC agent constructs a corresponding PoF and submits a transaction $tx' = P.SA.Withdrawn(amt, dest, PoF)$ to the parent subnet.
4. Upon ordering tx' , $P.SA.Withdrawn(amt, dest, PoF)$ verifies the PoF and transfers amt from SA (concretely, to src account representation within the SA) to $dest$ within the parent subnet.

Algorithm 7: Withdraw operation

```
1 ▶ owner of src:
2   └ submit  $tx = P/C.GWA.Withdraw(src, amt, dest)$ 
3 ▶  $P/C.GWA.Withdraw(src, amt, dest)$ :
4   └ deduct  $amt$  from  $src$  // "burns"  $amt$  in child
5 ▶ IPC agent:
6   └ upon  $tx = P/C.GWA.Withdraw(src, amt, dest)$  final at child do
7     └ create  $PoF(tx)$  // see section 7 for details
8     └ submit  $P.SA.Withdrawn(amt, dest, PoF)$ 
9 ▶  $P.SA.Withdrawn(amt, dest, PoF)$ :
10  └ verify( $PoF(tx')$ )
11  └ move  $amt$  coins from  $P.SA$  to  $dest$  // "unlocks"  $amt$  for  $dest$ 
```

4.4 Checkpointing

A checkpoint contains a representation of the state of the child subnet to be included in the parent subnet's replicated state. A checkpoint can be triggered by predefined events (e.g., periodically after a number of state updates, triggered by a specific user or set of users, etc.). A checkpoint is created as follows:

1. When the predefined checkpoint trigger is met (the IPC Agent, monitoring the child subnet's state, is configured with the checkpoint trigger), the IPC agent retrieves the corresponding checkpoint data ($chkp$) from the child subnet, along with the proof of its finality (PoF).
2. **[TODO:** Here the IPC Agent should decide (based on rights and some possible reward) whether to submit the Checkpoint transaction.]
3. The IPC agent submits a transaction $tx = P.SA.Checkpoint(chkp, PoF)$.
4. The $P.SA.Checkpoint(chkp, PoF)$ invocation, after verifying the PoF , includes $chkp$ in its state.

Algorithm 8: Checkpoint operation

```
1 ▶ IPC agent:
2   └ upon Checkpoint condition in child do
3     └  $chkp =$  obtain state snapshot from child
4     └ create  $PoF(chkp)$ 
5     └ submit  $P.SA.Checkpoint(chkp, PoF(chkp))$ 
6 ▶  $P.SA.Checkpoint(chkp, PoF(chkp))$ :
7   └ verify( $PoF(tx')$ )
8   └ save  $chkp$  in the state
9   └ [TODO: Expand on this. check if the checkpoint is the latest one and use a
      variable to store the latest checkpoint]
```

4.5 Propagating cross-net transactions

Unlike a "standard" transaction issued and submitted to a subnet by a user, a cross-net transaction is issued by the replicated logic of another subnet. Cross-net transactions are a means of interaction between smart contracts located on different subnets.

Since those smart contracts themselves are not processes (but mere parts of a subnet's replicated state), they cannot directly submit transactions to other subnets. IPC therefore provides a mechanism to propagate these transactions between subnets using a "*postBox*" and an IPC agent. In a nutshell, if a smart contract's logic produces a transaction for a different subnet, this transaction is saved the local Gateway Actor in a buffer that we call the *postBox*. The IPC agent, monitoring the *postBox*, then submits the transaction to the appropriate subnet.

Since, in general, we only rely on IPC Agents to be able to submit transactions to parents or children of a subnet whose state they observe, the IPC agent only propagates the transaction to the parent or child, depending on which is closer in the IPC hierarchy to the ultimate destination subnet. After such "one hop", the transaction is again placed in the *postBox* of the parent / child, and the process repeats until the transaction reaches its destination subnet.

The implementation of the Gateway Actor's *Propagate* function is sketched in Algorithm 9.

[**TODO:** Formalize subnet names and explain how the "*src*" is built.]

Algorithm 9: Cross-net transaction propagation functionality

```

1  ► GWA.Propagate(tx, src, dest, PoF):
2    verify(tx.PoF)
3    case dest = this subnet do
4      | apply tx
5    case dest requires going up the tree do
6      | postBox ← postBox ∪ (tx, S/src, dest)
7    case dest requires going down the tree do
8      | postBox ← postBox ∪ (tx, src/S, dest)
9  ► IPC agent:
10   upon new entry (tx, src, dest) in parent.GWA.postBox do
11     Create PoF proving that tx' has indeed been added to the list fo
12     cross-net transactions in the subnet
13     submit tx', augmented by PoF

```

4.6 Slashing

[**arp:** leave for later on with incentives and reconfiguration? it is hard to (meaningfully) talk about this without involving these concepts. Here a first attempt though:] Slashing is a penalty imposed on provably malicious validators. When validators of

a child subnet misbehave, other participants can report the misbehavior for these malicious validators to get punished (e.g. by losing a previously collateralized amount). Contrary to misbehaviors at a subnet with no parents, where misbehavers successfully perform their attack without escrow available, misbehaviors at the child can be resolved at the parent subnet, provided the misbehavers have not left the subnet. For this reason, a slash focuses on notifying the parent subnet as soon as possible, in the hope to stop an attempted attack. In particular, a slash on provably malicious validators of a child subnet is performed as follows:

- When the IPC agent of a correct participant identifies slashable misbehavior at subnet P/C from a set \mathcal{M} of malicious validators of subnet P/C , the IPC agent constructs a *Proof of Misbehavior* (PoM).
- The IPC agent then submits transaction $tx_P = P.SA.Slash(\mathcal{M}, PoM)$ at the parent and $tx_C = P/C.GWA.Slash.(\mathcal{M}, PoM)$ at the child.
- The parent subnet orders and executes tx_P . Once the parent's replicated state that includes tx_P becomes final, the IPC agent constructs a $PoF(tx_P)$ and submits a transaction $tx'_P = P/C.GWA.Slashed(\mathcal{M}, PoM, PoF(tx_P))$
- When the child's replicated state that includes tx_C becomes final (for some SMR-system-specific definition of finality that has been defined in the SA), the IPC agent constructs a corresponding $PoF(tx_C)$ and submits a transaction $tx'_C = P.SA.Slashed(\mathcal{M}, PoM, PoF)$ to the parent subnet.
- The parent subnet, upon ordering and executing either tx_P or tx'_C , penalizes the misbehavers and updates the state of SA .
- The child subnet, upon ordering and executing either tx_C or tx'_P , updates its state to reflect the penalization at the parent. [**arp:** This behavior at the child can later be updated to abort if it took place with tx_C and in fact tx'_P will never happen (i.e. attackers were faster and left subnet on time)]

Algorithm 10: Slash Functionality

```
1 [arp: this alg must be updated, but RSs to converge on text above first]
  input: -
2 ▶ Child SMR:
3   upon Proofs of fraud pofs generated do
4   |   Notify ⟨report, pofs⟩ to IPC agent
5 ▶ IPC agent:
6   upon ⟨report, pofs⟩ notified by child SMR do
7   |   Submit ⟨slash, pofs⟩ to parent SMR
8   upon [arp: State updated after slashing] do
9   |   [arp: Check child SMR rules are still satisfied, remedy/close otherwise?]
10 ▶ Parent SMR:
11   upon ⟨slash, pofs⟩ submitted by IPC agent do
12   |   Update SA state slashing/excluding participants Notify SA update to
      |   IPC agent
```

4.7 Removing a child subnet

A child subnet P/C can be removed from its parent P through a transaction invoking $P.GWA.RemoveChild(P/C)$. [mp: We will later define a mechanism to determine who has the right to do this and when.]

5 Incentives

In the previous sections, we defined the components of an IPC system and their roles in implementing the IPC functionality. Most of the functionality involved submitting transactions to subnets by an IPC Agent. However, in general, submitting transactions (and their subsequent execution by the subnet) is associated with a *cost* (often referred to as "gas"). We refer to the cost associated with a transaction as the *transaction fee*, measured in coins. An IPC Agent might need an incentive to participate in such a costly protocol.

Moreover, the replicas of a subnet might need to cooperate with IPC agents during the construction of Proofs of Finality. Unless non-cooperation can be detected and penalized (see ??), subnet replicas also might need an incentive to participate in the creation of a PoF.

This section describes mechanisms that can be used to incentivize participants running IPC agents to submit the required transactions and pay the corresponding transaction fees, as well as replicas to participate in PoF creation. It is *not* the goal of this section to provide a game-theoretic model of viable incentive mechanisms and their analyses. We merely present tools for implementing such mechanisms, to be used by those who design and implement concrete instances of IPC subnets.

5.1 Accounts

We assume that each IPC agent has accounts in both the parent and the child subnet and that the fees for the transaction the IPC agent submits to the respective subnets are deducted from the respective accounts. If the balance of the account is insufficient to pay the transaction fee, the transaction is considered invalid and is ignored by the subnet. We further assume that smart contracts (the *SA* and the *GWA*) can also hold funds and that their logic can distribute those funds among other accounts on their respective subnets.

Gateway Actor. The Gateway actor accumulates from its own subnet. For example, the subnet’s implementation can require a certain part of each transaction fee to be sent to the Gateway Actor.

Subnet Actor. The Subnet Actor accumulates funds from the subnet it governs. There are several ways how one can imagine the *SA* to be funded, for example, by withdrawals initiated by the child subnet’s *GWA*, or by periodically charging the child’s replicas for being included in the replica set.

5.2 Refunds and Rewards

An IPC Agent submitting a transaction that invokes a function of the *GWA* or the *SA* can receive a refund of the transaction fee directly from the invoked smart contract, according to arbitrary rules defined in the smart contract’s logic. For example, such a rule could be to credit the submitting IPC Actor’s account with the amount of the transaction fee, augmented by a fixed (or proportional) reward.

To incentivize the replicas of a subnet to collaborate with the IPC agent on the creation of Proofs of Finality, a similar mechanism can be deployed. For example, a valid PoF would include metadata, where the replicas that participated in its creation could insert an address to receive a reward when the PoF is accepted.

6 An Instance of IPC

Here we describe the particular choices implemented by the Consensus Lab team for the reference implementation of IPC. The current implementation considers Filecoin as the root subnet, and Trantor as child subnets. For our interest it is important to note that Trantor is a BFT consensus protocol with immediate finality, and Filecoin is a longest chain style protocol with probabilistic finality.

The two main conceptual choices in the implementation are: (i) batching upward transactions – both from the same type as well as batching different

kinds of transactions together (which includes withdrawals, checkpointing and postBox transactions that must traverse the parent); (ii) using multisigs for verifying finality of a child state, and local finality check for the finality of a parent state.

Batching through the *GWA*. All upward transactions are made via the *GWA* of the child. The transactions are batched there until it is time to checkpoint at the parent (specifically, every Δ child blocks). This way the *GWA* also serves as the single data structure (queue) storing the upward transactions, therefore, an IPC agent only checks this single place to get all necessary info from the child subnet. When the checkpoint batch is executed at the parent, it is done atomically. Since postBox transactions are included in the batch, the atomic execution of the batch also depends on the postBox functionality at the parent handling those transactions (recall that the parent’s postBox lays in the parent’s *GWA*). Therefore, *SA* first commits the cid of the checkpoint batch without executing the included transactions, and then triggers the *GWA* to examine the batch as well. The *GWA* checks whether there are sufficient funds to handle the postBox transactions in the batch. If the funds are insufficient, the entire checkpoint batch fails. If the *GWA* approves the batch, then it is executed — both at *SA* and at *GWA*.⁶

Checking for Finality. Recall that we use Trantor as the consensus engine for child subnets. Therefore, a child’s transaction (or state) *tx* is accepted at the parent as final by providing a *PoF* containing 2/3 of the child’s validators signatures⁷ on *tx*. In other words, given that *tx* is a cid, the parent’s subnet call to *SA.verifyGlobalFinality(tx, PoF)* returns True iff *PoF* contains signatures (on *tx*) of validators with at least 2/3 of the voting rights in the child. The voting rights are measured according to what is written in *SA* for the epoch containing *tx*.

For considering a transaction *tx* at the parent as final, we use the fact that a participant has view into a version of the parent subnet (through its local parent replica process). In this case, the *PoF* contains the block height *h* (and pointer to that block) at the parent subnet. A child replica then asserts with its parent that the state is final by checking with its local version of the parent blockchain at height *h*. If the local version at the parent

⁶Batching the postBox transactions together with the checkpoint and withdrawals, as well as entangling the *GWA* functionality (postBox transactions) with the *SA* functionality (withdrawals and checkpoint transactions) in a single atomic execution are design choices. These choices benefited the development velocity by having a single mechanism to handle everything and by transferring the responsibility of re-transmitting failed postBox transactions to the validators rather than it being the responsibility of *SA* (which requires logic implementation).[\[gg: Alfonso, please explain...\]](#)

⁷The current implementation relies on collecting multiple signatures. A next step in the implementation road-map is to offer a threshold signature mechanism instead of using a multisig. For now, multisigs serve the purpose of an MVP implementation.

replica did not reach height h yet, the child replica considers the state to currently be non-final/non-valid. The child replica checks again when the parent replica reaches height h . [TODO: Verify with REs][gg: If it is not done, then it should be done.]

IPC agent and metadata. In the previous sections we considered that every pair of parent and child will have an independent IPC agent process. In fact, the implementation manages to execute one single IPC agent for the entire tree of subnets that may be of relevance to the participant. This process can be executed either as a daemon or as a CLI. In the latter, though, the IPC agent can participate in neither checkpointing nor propagating cross-net transactions. Additionally, the reference implementation makes use of IPFS content addressable data, in that data is stored where relevant and referred to with a content identifier (CID) elsewhere.

Compressed accounting. In section 3 we mentioned that *SA.accounts* can include fine grained accounting data. However, in the current implementation *SA.accounts* contain a single variable representing the sum of all the individually locked funds at the parent which are dedicated to the child subnet (we call this value *circulating supply*). This has the obvious benefit of reducing the space complexity of *SA*.

Checkpointing. We show in Algorithm algorithm 11 the main design choices made by the reference implementation. A checkpoint is first triggered every *Delta* blocks decided at the child subnet. If the latest block decided meets this condition, and if the participant’s child subnet is a validator according to the state stored at the parent, then the IPC agent starts computing the checkpoint as follows:

- The IPC agent obtains a state snapshot from the child’s subnet.
- The IPC agent obtains the CIDs of all new grandchildren’s checkpoints, and of the upward postBox messages
- The IPC agent computes the checkpoint *chkp*, compressing its state by computing only the additional changes from the latest checkpoint stored at the parent’s SA, and creates a PoF (*chkp*)
- The IPC agent submits $tx = P.SA.Checkpoint(chkp, PoF(chkp))$
- The parent subnet, upon ordering and executing tx , saves *chkp* in the state. [TODO: Saves a cid after submitting it to IPFS? or the actual checkpoint?] [TODO: and updates also *GWA* with the new postbox messages?]

Algorithm 11: Checkpoints IPC reference implementation

```

1  ► IPC agent:
2  │   upon newBlock from child subnet do
3  │   │   if newBlock.blockheight mod  $\Delta = 0$  then
4  │   │   │   if P.SA.isValidator(Self) then
5  │   │   │   │   state  $\leftarrow$  obtain state snapshot from child;
6  │   │   │   │   gcChkps  $\leftarrow$  query child's GWA for grandchildren's checkpoints
7  │   │   │   │   postboxmsgs  $\leftarrow$  query child's GWA for upward postbox messages
8  │   │   │   │   pChkps  $\leftarrow$  query parent's GWA for previous checkpoints
9  │   │   │   │   chkp  $\leftarrow$  createChkp(state, pChkps.latestChkp, gcChkps,
10 │   │   │   │   │   postboxmsgs) // compress state
11 │   │   │   │   create PoF(chkp) // multisig
12 │   │   │   │   submit P.SA.Checkpoint(chkp, PoF(chkp))
13 │   │   │   └─┘
14 │   │   └─┘
15 │   └─┘
16 ► P.SA.Checkpoint(chkp, PoF(chkp)):
17 │   verify(PoF(tx'))
18 │   save chkp in the state
19 │   [TODO: Call GWA to store there the upward postbox msgs]

```

Incentives At the moment, validators get rewarded for executing the checkpoint algorithm by charging an IPC fee on all transactions traversing the postbox. This incentivises validators in participating on the checkpointing functionality, even if that costs them a fee to be paid for the transaction at the parent. [TODO: TBC. No governance account, but participation incentives by:

1. Additional IPC fee for validators to relay
 2. Incentivizing checkpoint submission as a result of batching with other crossnet messages (which contain the IPC fees), further justifying checkpoints being stored at GW.
 3. Keeping Stake at SA and slashing through fraud proofs
-]

7 Verifying the Finality of tx

[arp: I think this section should contain much more than this (but that perhaps this section does not follow our timelines for document completion (more of a complement of the document). Particular content here imo: Analysis for improvements wrt reference implementations (i.e. threshold signatures instead of everyone submitting checkpoints, governance account instead of no incentives, etc.); and Comprehensive list of different approaches for functionality/functions (like we had in the legacy document).][gg: Agreed] A main ingredient in any Interplanetary Consensus implementation is the creation and verification of a finality proof for a given tx in some subnet. In the previous sections we left these functions opaque. For example, *SA.verifyGlobalFinality*(tx, PoF) was used by the parent replica to verify the finality at the child subnet of tx . The creation of *PoF* and the verification method at the child replica (for transactions of that occur at

the parent subnet), are only hinted by plain text. There are multiple ways to implement these functionalities, each with its own trade-offs. Below we propose several such implementations.

References

- [1] IPC Glossary. <https://docs.google.com/document/d/15pA7ahjeA-HY0l8Pxj0n6PxEswYlRVrZ112MJuRR0fY/edit?usp=sharing>.

A Glossary

[**TODO:** (Marko)Add IPC Glossary [1] here and move most of model down here]